

# EXPERIMENT 4

OBJECTIVE : WAP to evaluate the performance of implemented three-layer neural network with variations in activation functions, size of hidden layer, learning rate, batch size and number of epochs.

DESCRIPTION OF MODEL(Configuration : Hidden layer , Learning rate)

The model is a fully connected feed-forward neural network designed to classify handwritten digits from the MNIST dataset.

Model follows a structured approach using forward propagation, loss calculation, backpropagation, and optimization to learn the correct classification.

## 1. Network Architecture 1

- Input layer : Each MNIST image is  $28 \times 28$  pixels (784 pixels)--> 784 input size
- Hidden Layer 1 : 256 neurons , Activation Function: ReLU
- Output Layer: 10 neurons (For each digit from 0 to 9) , No activation because it will be handled by softmax in the loss function.

Model parameters(hyperparameters)

- Number of epochs : 50
- Learning rate : 0.01
- Batch size : 10
- Loss function : Softmax Cross-Entropy/Categorical cross-entropy
- Optimiser : Adam

## 2. Network Architecture 2 (For 15 configurations)

- Input layer : Each MNIST image is  $28 \times 28$  pixels (784 pixels)--> 784 input size
- Hidden Layers (2) : Neuron Configurations- [(160, 100), (100, 100), (100, 160), (60, 60), (100, 60)] , Activation Function: ReLU
- Output Layer: 10 neurons (For each digit from 0 to 9) , No activation because it will be handled by softmax in the loss function.

Model parameters(hyperparameters)

- Number of epochs : 50
- Learning rates : [ 0.01 , 0.1 , 1.0]

- Batch size : 10
- Loss function : Softmax Cross-Entropy/Categorical cross-entropy
- Optimiser : Adam

## PYTHON IMPLEMENTATION

```
import tensorflow as tf

import tensorflow_datasets as tfds

import numpy as np

import matplotlib.pyplot as plt

import time

from sklearn.metrics import confusion_matrix

import seaborn as sns


# Load and preprocess the MNIST dataset

def preprocess(image, label):

    image = tf.cast(image, tf.float32) / 255.0 # Normalize to [0,1]

    image = tf.reshape(image, [-1]) # Flatten to 784

    label = tf.one_hot(label, depth=10) # Convert to one-hot encoding

    return image, label


# Load dataset and apply preprocessing

mnist_dataset = tfds.load("mnist", split=["train", "test"], as_supervised=True)

train_data = mnist_dataset[0].map(preprocess).shuffle(10000).batch(10)

test_data = mnist_dataset[1].map(preprocess).batch(10)
```

```
# Define neural network parameters
```

```
input_size = 784
```

```
hidden_layer1_size = 160
```

```
hidden_layer2_size = 100
```

```
output_size = 10
```

```
learning_rate = 0.01
```

```
epochs = 50
```

```
# Initialize weights and biases
```

```
W1 = tf.Variable(tf.random.normal([input_size, hidden_layer1_size]))
```

```
b1 = tf.Variable(tf.zeros([hidden_layer1_size]))
```

```
W2 = tf.Variable(tf.random.normal([hidden_layer1_size, hidden_layer2_size]))
```

```
b2 = tf.Variable(tf.zeros([hidden_layer2_size]))
```

```
W_out = tf.Variable(tf.random.normal([hidden_layer2_size, output_size]))
```

```
b_out = tf.Variable(tf.zeros([output_size]))
```

```
# Forward pass function
```

```
def forward_pass(X):
```

```
    hidden_layer1 = tf.nn.relu(tf.matmul(X, W1) + b1)
```

```
    hidden_layer2 = tf.nn.relu(tf.matmul(hidden_layer1, W2) + b2)
```

```
    output_layer = tf.matmul(hidden_layer2, W_out) + b_out # No activation (logits)
```

```
    return output_layer
```

```
# Define loss function
```

```
def loss_fn(logits, labels):
```

```
    return tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=logits,
labels=labels))
```

```
# Define optimizer
```

```
optimizer = tf.optimizers.Adam(learning_rate)
```

```
# Training step function
```

```
def train_step(X, Y):
```

```
    with tf.GradientTape() as tape:
```

```
        logits = forward_pass(X)
```

```
        loss = loss_fn(logits, Y)
```

```
    gradients = tape.gradient(loss, [W1, b1, W2, b2, W_out, b_out])
```

```
    optimizer.apply_gradients(zip(gradients, [W1, b1, W2, b2, W_out, b_out]))
```

```
    return loss
```

```
# Compute accuracy
```

```
def compute_accuracy(dataset):
```

```
    total_correct = 0
```

```
    total_samples = 0
```

for X, Y in dataset:

logits = forward\_pass(X)

correct\_pred = tf.equal(tf.argmax(logits, 1), tf.argmax(Y, 1))

total\_correct += tf.reduce\_sum(tf.cast(correct\_pred, tf.float32))

total\_samples += X.shape[0]

return total\_correct / total\_samples

# Track loss and accuracy

loss\_history = []

accuracy\_history = []

start\_time = time.time()

# Training loop

for epoch in range(epochs):

avg\_loss = 0

total\_batches = 0

for batch\_x, batch\_y in train\_data:

loss = train\_step(batch\_x, batch\_y)

avg\_loss += loss

total\_batches += 1

```
avg_loss /= total_batches

train_acc = compute_accuracy(train_data)

loss_history.append(avg_loss.numpy())

accuracy_history.append(train_acc.numpy())

print(f"Epoch {epoch+1}, Loss: {avg_loss:.4f}, Training Accuracy: {train_acc:.4f}")

end_time = time.time()

execution_time = end_time - start_time


# Test the model

test_acc = compute_accuracy(test_data)

print(f"Test Accuracy: {test_acc:.4f}")


# Plot Loss Curve

plt.figure(figsize=(10, 5))

plt.plot(loss_history, label='Loss')

plt.xlabel('Epochs')

plt.ylabel('Loss')

plt.title('Loss Curve')

plt.legend()

plt.show()


# Plot Accuracy Curve
```

```
plt.figure(figsize=(10, 5))

plt.plot(accuracy_history, label='Accuracy', color='orange')

plt.xlabel('Epochs')

plt.ylabel('Accuracy')

plt.title('Accuracy Curve')

plt.legend()

plt.show()
```

```
# Compute confusion matrix
```

```
y_true, y_pred = [], []
```

```
for X, Y in test_data:
```

```
    logits = forward_pass(X)
```

```
    predictions = tf.argmax(logits, axis=1).numpy()
```

```
    labels = tf.argmax(Y, axis=1).numpy()
```

```
    y_pred.extend(predictions)
```

```
    y_true.extend(labels)
```

```
conf_matrix = confusion_matrix(y_true, y_pred)
```

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=range(10),
yticklabels=range(10))
```

```
plt.xlabel("Predicted")
```

```
plt.ylabel("Actual")
```

```
plt.title("Confusion Matrix")
```

```
plt.show()
```

```
print(f"Execution Time: {execution_time:.2f} seconds")
```

## DESCRIPTION OF CODE

- tensorflow – Used for building and training the neural network.
- tensorflow\_datasets – Used to load the MNIST dataset.
- numpy – Used for numerical operations and working on arrays.
- matplotlib.pyplot – Used to plot loss and accuracy curves.
- time – Used to measure the execution time over training loop.
- sklearn.metrics.confusion\_matrix – Used to compute the confusion matrix for performance evaluation.
- seaborn – Used to visualise the confusion matrix as heatmap.
- First we will load the mnist dataset.

preprocess() :

- Dataset is loaded and preprocessing is done . Pixel values are normalised between [0,1] .
- Dataset is converted to a 1D array (784).
- One hot encoding is done to convert categorical labels to vector having binary values , e.g. class label 2 –[0 ,0, 1, 0, 0, 0, 0, 0, 0, 0].
- mnist dataset is split into train and test dataset each containing tuple (image , label).
- map(preprocess) : Applies preprocessing to every image in the dataset.
- shuffle(10000) : Randomly shuffles 10,000 images to prevent model learn by order of images.
- batch(10) : Divides dataset into mini-batches of size 10 images for training. For 10 forward passes , backward propagation will be performed once.
- Neural network parameters are defined.

Weights and bias initialisation :

- Initialised using tf.Variable . Weights (W1, W2, W\_out) : Randomly initialized
- Biases (b1, b2, b\_out) : Initialized as zero

forward\_pass() :

- Layer 1 output :  $A1 = \text{sigmoid}(XW1 + b1)$
- Layer 2 output :  $A2 = \text{sigmoid}(A1W2 + b2)$
- Output Layer (Logits):  $\text{Output}(Z) = (A2W_{\text{out}} + b_{\text{out}})$
- Returns logit values.



loss\_fn() :

- Softmax Cross-Entropy Loss: Measures how different the predictions are from true labels.
- $\text{Loss} = -\sum (Y_i) \log(\text{softmax}(Z_i))$
- `tf.reduce_mean()`: Computes average loss within a batch.
- Adam Optimizer : Optimizer adjusts the weights and biases to minimize the loss function during training.

train\_step() :

- `tf.GradientTape()`: Help in resource management . Stores sequence of operations (like matrix multiplication , logits , loss values ,which parameter contributed more in loss) inside the block and calculate gradients automatically.
- `tape.gradients()` calculate gradient of loss with respect to each weight and bias . Memory is freed up as soon as it is called.
- `optimizer.apply_gradients()`: Updates weights using the calculated gradients , zip pairs gradient with respective parameter.

compute\_accuracy() :

- Helps calculate training and test accuracy (correct predictions/total predictions).
- Prediction - (argmax): Returns the index of the highest probability class.
- `equal()`: Compares predictions with true labels.

Training loop:

- Loops over 50 epochs.
- Calculates the loss and training accuracy . Prints loss and accuracy value for each epoch . These values are added in loss history and accuracy history lists.
- Execution time is calculated over the training loop.
- Test accuracy is calculated.
- Loss curve , Accuracy curve and Confusion matrix are plotted.

## PERFORMANCE EVALUATION

- Performance has been evaluated by plotting Loss curve , Training Accuracy curve and Confusion matrix.
- The training and test accuracy for **Network architecture 1** is Training Accuracy: 33.64% , Test Accuracy: **33.43%** as it just contains a single hidden layer and batch-size is also small(10) , it has executed in **least time(10308.36 seconds)**.
- The training and test accuracy for Network architecture 2 having having 15 configurations vary according to the model parameters.
- The training and test accuracy achieved is less for **learning rate 0.01** but maximum train and test accuracy reached are **Training Accuracy: 91.76%** , **Test Accuracy: 91.18%** for configuration of **hidden layer (60,60)** .
- For **learning rate values 0.1 and 1.0** the training and test accuracy were nearly same and have just reached **9 – 10%** .
- **Maximum execution time** required (15 combinations) was for configuration having learning rate 1.0 (**14831.84 seconds**) for configuration of **hidden layer (60,60)**.
- **Minimum execution time** required (15 combinations) was for configuration having learning rate 0.01 (**12698.1 seconds**) for configuration of **hidden layer(60,60)**.
- The **execution time** required is higher for learning rate 1.0 than for learning rate values 0.01 and 0.1 .
- Execution time ranges from approximately 12000 to 15000 seconds.

## My Comments(Limitations and Scope for Improvement)

- Maximum train and test accuracy reached are **Training Accuracy: 91.76%** , **Test Accuracy: 91.18%** for configuration of **hidden layer (60,60)** .
- The accuracies achieved are very less in maximum cases which is just about **9-10%**.
- The execution time is high.
- To improve accuracy and decrease execution time , a different activation function can be used , batch size can be increased , the learning rate value can be changed , the number of hidden layers and number of hidden layer neurons need to be changed .